

# Technical Architecture: Centralized Event Review Platform

**\*\*Version:\*\*** 1.1

**\*\*Date:\*\*** May 6, 2025

**\*\*Status:\*\*** Proposed

## 1. Introduction

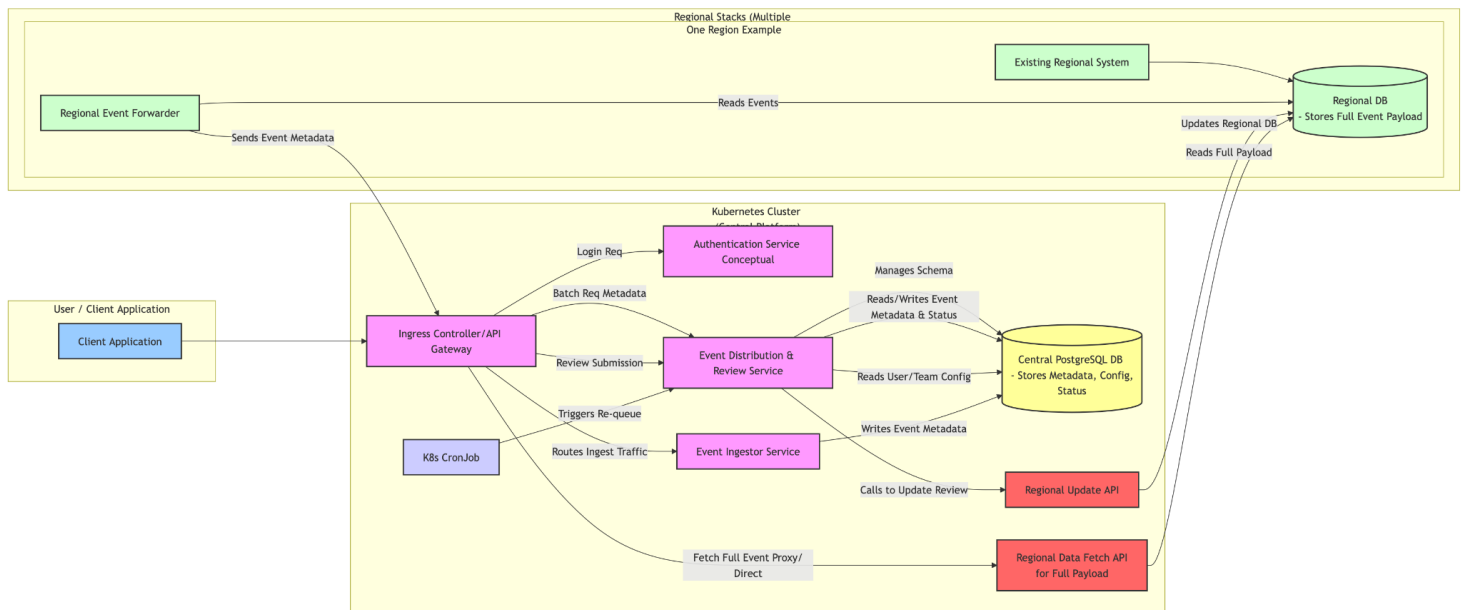
This document outlines the technical architecture for the Centralized Event Review Platform. The platform aims to provide internal engineering teams with a unified interface to review and rate events generated by edge devices deployed across multiple regional AWS stacks (e.g., US, CA, APAC, EU). A key design consideration is to meet regional data compliance requirements by ensuring sensitive regional data remains within its originating region, while still enabling a centralized review workflow.

## 2. Core Architectural Principles

- **Microservice-Oriented:** The central platform is composed of distinct, independently deployable microservices to promote scalability, maintainability, and technology flexibility.
- **Centralized Workflow, Decentralized Data:** Core review workflow management (assignment, status tracking) is centralized. However, the full, potentially sensitive, event payloads remain in their originating regional datastores. The central platform primarily handles metadata.
- **API-Driven:** Communication between services, regional forwarders, and the client application is managed via well-defined APIs.
- **Containerization & Orchestration:** Services are designed to be containerized (Docker) and orchestrated by Kubernetes for scalability, resilience, and consistent deployments across environments.
- **Data Compliance by Design:** The architecture prioritizes keeping sensitive regional data within its designated region, with the central platform accessing or referencing it as needed without persistent central storage of the full payload.

### 3. High-Level Architecture Diagram

The following diagram illustrates the major components and their interactions:



### 4. Component Descriptions

#### 4.1. Central Platform Components

- **Client Application:** The frontend through which users log in, view assigned event metadata, request full event details (via the platform), and submit their reviews (ratings, comments, approval/rejection).
- **Ingress Controller / API Gateway (e.g., Nginx Ingress):**
  - Single entry point for all external traffic (from clients and regional forwarders).
  - Handles SSL termination, routing requests to appropriate backend services based on path/host.
  - Can enforce policies like rate limiting, authentication at the edge.
- **Authentication Service (Conceptual):**
  - Responsible for user authentication and session management (e.g., issuing JWTs).
  - (For POC, this is stubbed/bypassed, with user identity passed via headers).

- **Event Ingestor Service:**
  - Provides an API endpoint (e.g., /api/v1/internal-ingest/event) for Regional Event Forwarders.
  - Receives event *metadata* from various regions.
  - Validates incoming metadata.
  - Stores the metadata in the Central Data Store with an initial 'Pending' status and region of origin.
- **Event Distribution & Review Service:**
  - The core workflow engine.
  - Handles user requests for event batches (metadata).
  - Assigns unique batches to users based on team/region configurations, ensuring no concurrent assignments using database-level locking.
  - Processes review submissions:
    - Calls the appropriate Regional Update API to record the review outcome in the regional datastore.
    - Updates the event's status in the Central Data Store to 'Completed' (or manages deletion for stricter compliance).
  - Exposes an internal endpoint (e.g., /api/v1/internal/trigger-requeue) for the CronJob.
  - Manages the database schema for the Central Data Store via migrations (node-pg-migrate).
- **Central Data Store (PostgreSQL):**
  - Stores event *metadata* (not full payloads), including event\_id (central UUID), external\_event\_id, region\_code, status ('Pending', 'Assigned', 'Completed'), assignment details (assigned\_user\_id, assigned\_at), and review outcomes (review\_user\_id, reviewed\_at, review\_decision, review\_comment).
  - Stores user profiles, team configurations, and team-to-region mappings.
- **Kubernetes CronJob:**
  - Periodically calls an internal endpoint on the Event Distribution & Review Service to trigger the re-queuing logic for events that have timed out or whose assigned user session has expired.

## 4.2. Regional Stack Components

- **Existing Regional System:** The source system generating events.
- **Regional Database:** Stores the full, original event payloads, adhering to regional data residency and compliance requirements.
- **Regional Event Forwarder:**
  - A lightweight component/agent deployed in each region.
  - Monitors the Regional Database or event pipeline for new events designated for review.
  - Extracts only the necessary *metadata* (non-sensitive summary or reference).
  - Securely sends this metadata to the central Event Ingestor Service via the central

Ingress.

- **Regional Update API:**
  - An API endpoint exposed by each regional stack.
  - Called by the central Event Distribution & Review Service to write back review outcomes (approval/rejection, ratings, comments) to the original event record in the Regional Database.
- **Regional Data Fetch API:**
  - An API endpoint exposed by each regional stack.
  - Called by the Client Application (likely proxied through the central Ingress) to retrieve the *full event payload* for display to the user during the review process, using metadata like `external_event_id` and `region_code`.

## 5. Data Flow Summary

1. **Ingestion:** Regional Forwarder sends event METADATA via Ingress to the central Ingress Service.
2. **Storage (Central):** Ingestor Service stores this METADATA in the Central DB with 'Pending' status.
3. **Assignment Request:** Client App requests a batch of events (METADATA) from the Distribution Service (via Ingress).
4. **Assignment Logic:** Distribution Service queries Central DB for eligible 'Pending' metadata, assigns a batch to the user, and updates status to 'Assigned' in Central DB.
5. **Full Data Fetch (Review Time):** For review, the Client App (or Central Platform as proxy via Ingress) requests the FULL PAYLOAD from the appropriate Regional Data Fetch API using metadata (e.g., `external_event_id`, `region_code`).
6. **Review Submission:** Client App submits the review outcome to the Distribution Service (via Ingress).
7. **Write-Back (Regional):** Distribution Service calls the relevant Regional Update API to store the review outcome in the Regional DB.
8. **Finalize (Central):** Distribution Service updates the METADATA status in Central DB (e.g., to 'Completed' or deletes the metadata record for stricter compliance).
9. **Re-queue:** The K8s CronJob periodically triggers the Distribution Service to check for timed-out 'Assigned' events and revert their status to 'Pending' in the Central DB.

## 6. Data Compliance Strategy

- **Regional Data Residency:** Full, sensitive event payloads are stored and remain within their originating regional datastores.
- **Central Metadata Store:** The central platform primarily stores and processes non-sensitive metadata or references necessary for the review workflow. The `event_payload` field in the central events table is intended for this metadata, not the full regional payload.

- **On-Demand Full Payload Fetch:** Full event details for review are fetched directly from the regional source when needed by the user, minimizing central storage of sensitive data.
- **Secure Regional Write-Back:** Review outcomes are securely transmitted back to the originating regional datastore.
- **Central Record Management Post-Review:** After successful regional write-back, central metadata records are marked 'Completed'. For enhanced compliance, these records can be configured for archival and eventual purging, or even immediate deletion, based on retention policies and audit requirements.

## 7. Key Technology Choices (POC)

- **Backend Services:** Node.js with Express.js and TypeScript.
- **Central Database:** PostgreSQL.
- **Database Migrations:** node-pg-migrate.
- **Containerization:** Docker.
- **Orchestration:** Kubernetes (Minikube for local POC).
- **Ingress Controller:** Nginx Ingress Controller (Conceptual).

## 8. Scalability & Maintainability

- **Microservices:** Allow independent scaling and development of components.
- **Kubernetes:** Provides auto-scaling, self-healing, and rolling updates.
- **Stateless Services:** Application services are designed to be stateless where possible, relying on the Central DB for state.
- **Database:** PostgreSQL can be scaled (e.g., read replicas, connection pooling) as needed.
- **Asynchronous Processing (Future):** For very high ingestion rates, message queues could be introduced between forwarders and the ingestor, or for other background tasks.

## 9. Future Considerations

- **Full Authentication & Authorization:** Implement robust OAuth2/OIDC-based authentication and fine-grained role-based access control (RBAC).
- **Enhanced Monitoring & Alerting:** Integrate comprehensive logging, metrics (Prometheus, Grafana), and tracing (OpenTelemetry, Jaeger/Zipkin).
- **Client Application:** Develop the full-featured client application for reviewers.
- **Advanced Regional API Integration:** Standardize and secure regional APIs further.
- **Data Archival/Purging Strategy:** Formalize policies for central metadata retention.
- **Configuration Management:** For production, use more sophisticated configuration management tools